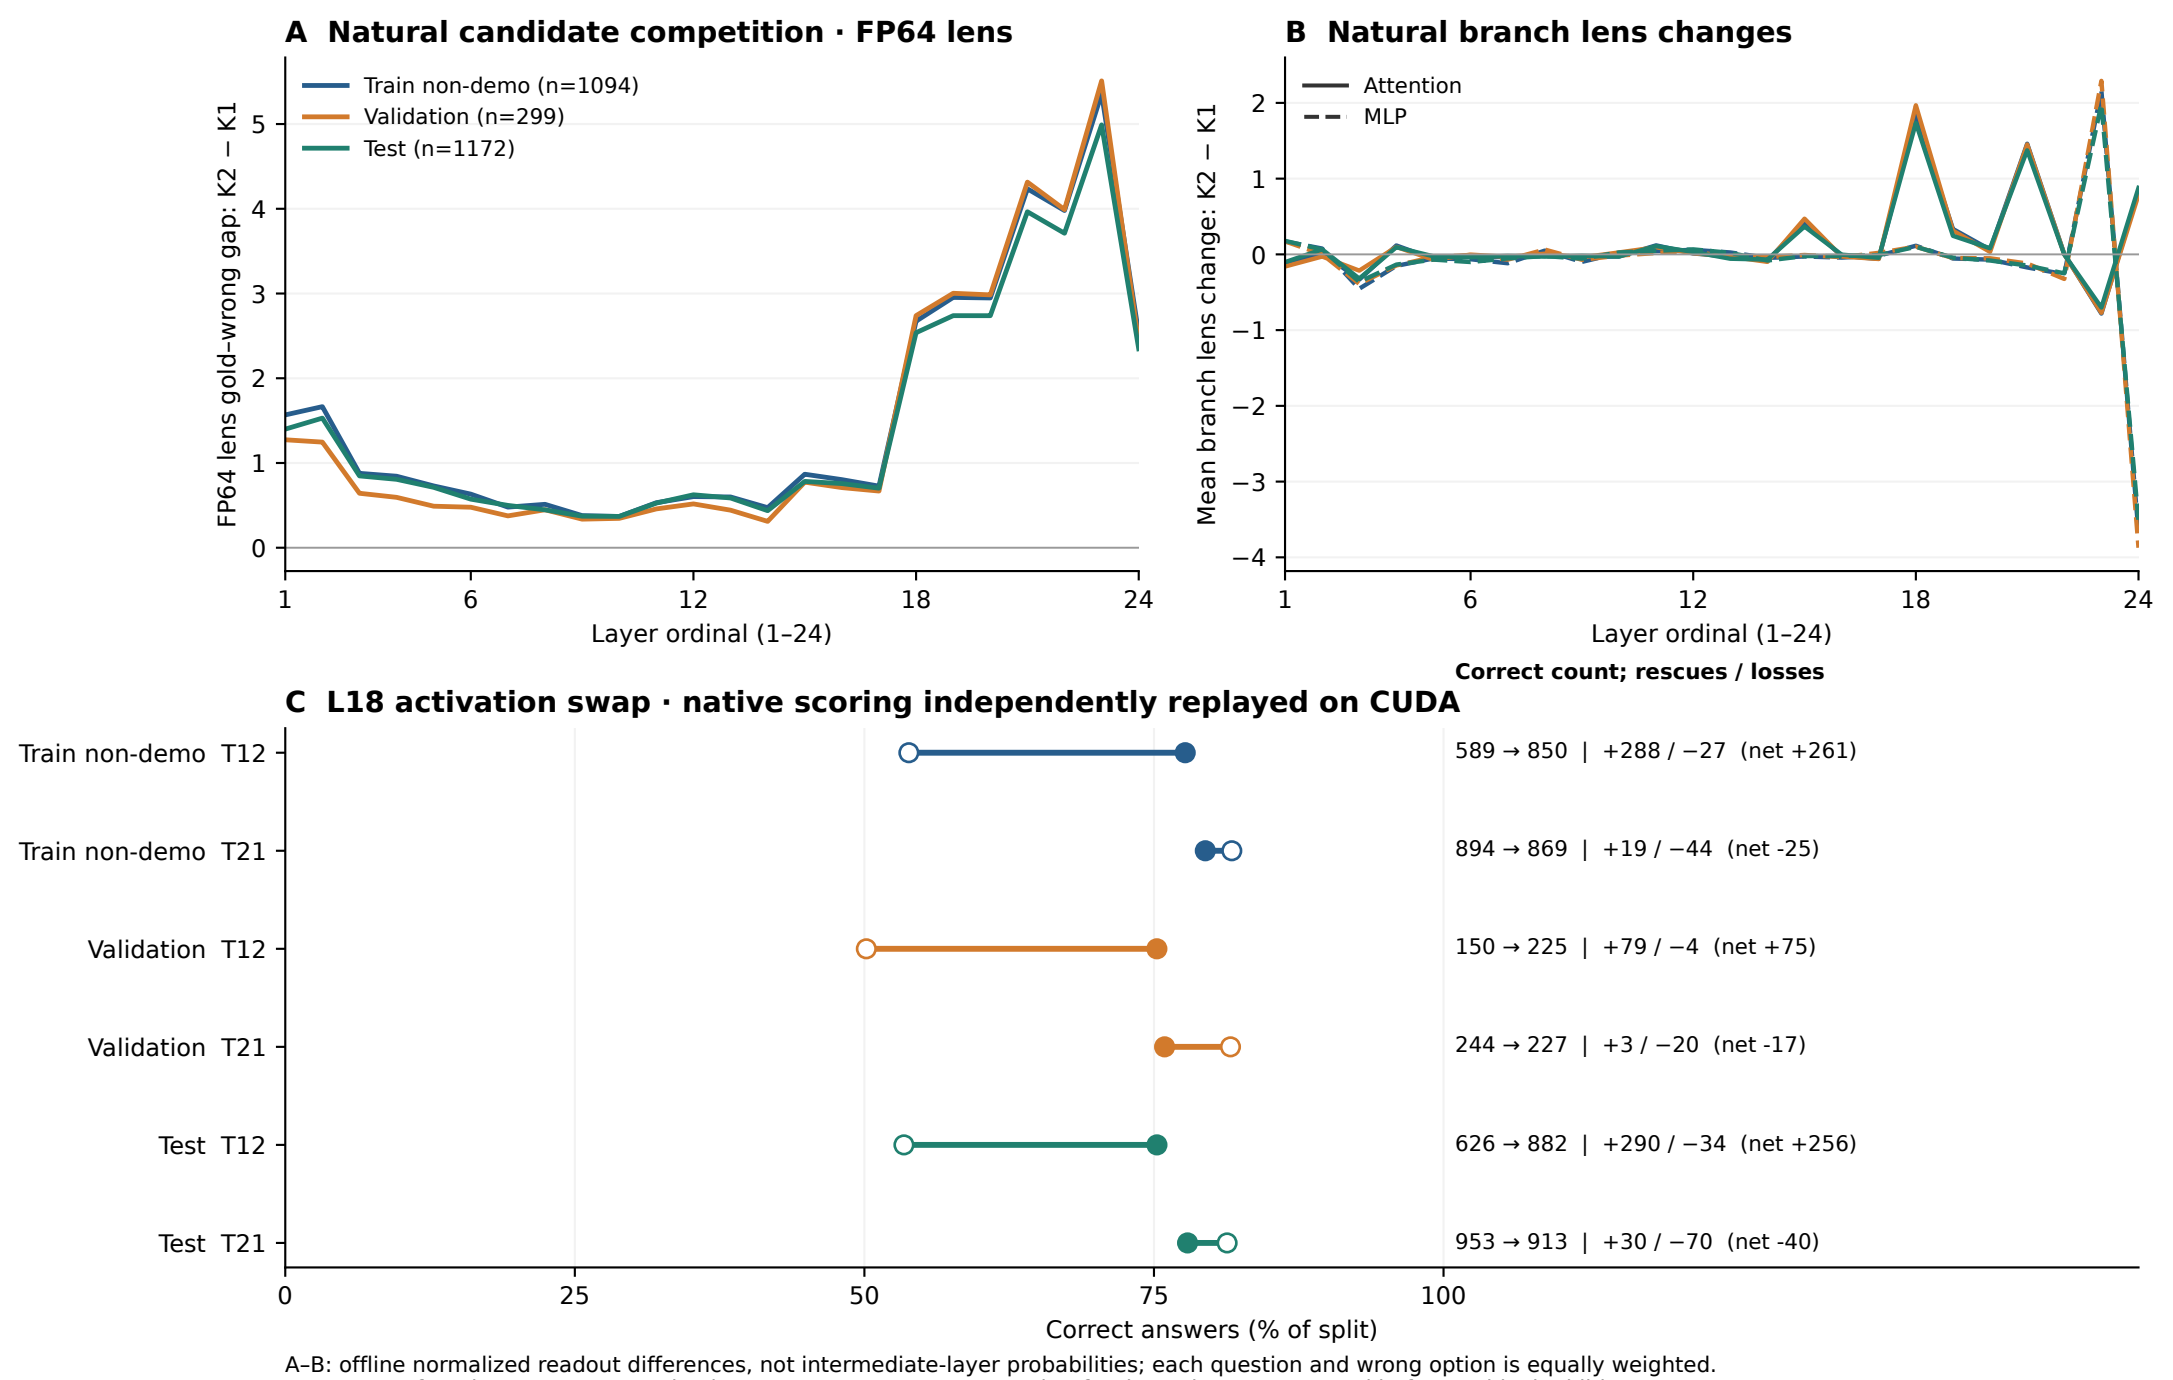

A–B: offline normalized readout differences, not intermediate-layer probabilities; each question and wrong option is equally weighted.
C: T12 transfers the K2 L18 R attention increment to K1; T21 reverses it, after input_layernorm_2 and before residual addition.
Open = baseline; filled = swap. Displayed layer ordinal: 1–24. Train excludes demonstrations.